

HOMEWORK 1

Question 2.1

Describe a situation or problem from your job, everyday life, current events, etc., for which a classification model would be appropriate. List some (up to 5) predictors that you might use.

I'm currently working on a project to forecast traffic utilization for a telecom company and could apply a classification model to differentiate traffic patterns. Predictors that could be used include:

1. Time-series predictor – Auto-regressive integrated moving average (ARIMA) models have been used by researchers to predict yearly growth in network traffic
2. Neural Networks-Based Predictors – Deep learning algorithms have been used for traffic prediction in wireless mesh networks
3. Wavelet Transform-Based Predictors – Used for multi-scale prediction as the naturally transforms the network wavelet into multiple resolutions depending on the size and range of the wave

Question 2.2

1. Using the support vector machine function `ksvm` contained in the R package `kernlab`, find a good classifier for this data. Show the equation of your classifier, and how well it classifies the data points in the full data set.

Multiple test cases were run on the `ksvm` model using C values ranging from 0.001 to 10000 but did not result in a better training error score. The model returns a training error of : 0.136086 which signifies the probability of prediction errors. The equation of the classifier is: $-0.0005(A1) - 0.01(A2) - 0.008(A3) + 0.01(A8) + 0.5(A9) - 0.002(A10) - 0.001(A11) + 0.0003(A12) - 0.2(A14) + 0.06(A15) - 40.4 = 0$

Note: I created multiple SVM models in Python using the `sklearn` library and was able to produce an accuracy score of 99.4% using the Radial basis function kernel. The code is included below the R solution.

```
> d0 = read.table('credit_card_data-headers.txt', header=TRUE)
> d0[,11] = replace(d0[,11], d0[,11] == 0, -1)
> d0[,11] = as.factor(d0[,11])
> colnames(d0)[11] = 'class'
> predictors = as.matrix(d0[,1:10]); classes = as.matrix(d0[,11])
>
> library(kernlab)
> model = ksvm(predictors, classes, type="C-svc", kernel="vanilladot", C=0.2, scaled=TRUE)
Setting default kernel parameters
> model
Support Vector Machine object of class "ksvm"
```

SV type: C-svc (classification)
parameter : cost C = 0.2

Linear (vanilla) kernel function.

Number of Support Vectors : 196

Objective Function Value : -36.2851

Training error : 0.136086

#Calculate the weights & constant

```
> w = colSums(predictors[model1@SVindex,1:10] * model1@coef[[1]])
```

```
> w
```

A1	A2	A3	A8	A9
-5.115191e-04	-9.931841e-03	-8.163952e-03	1.086537e-02	5.013006e-01
A10	A11	A12	A14	A15
-1.628945e-03	-1.024945e-03	-2.559994e-04	-2.016443e-01	5.587158e+02

```
> a0 <- sum(w*predictors[1,1:10]) - model1@b
```

```
> a0
```

```
[1] -40.44358
```

Python Solution using the same data:

```
from sklearn.svm import SVC
import pandas as pd
import numpy as np
```

```
#Reading in the data
```

```
file_loc = (r'C:\Users\Mitchell.Ramey\Documents\credit_card_data-headers.txt')
```

```
file = open(file_loc)
```

```
file = file.read()
```

```
df = pd.read_csv(file_loc, delim_whitespace=True)
```

```
#Seperating the Classifier/Predictors
```

```
clssfr = df.iloc[:,10:]
```

```
clssfr2 = df.loc[:, 'R1']
```

```
predictors = df.loc[:, 'A1':'A15']
```

```
#Fitting the model
```

```
clf = SVC(gamma='scale', kernel = 'linear')
```

```
clf_rbf_auto = SVC(gamma = 'auto')
```

```
clf_rbf_scale = SVC(gamma = 'scale')
```

```
Y = np.array(clssfr)
```

```
Y = np.reshape(Y, (654,))
```

```
y = Y
```

```

X = np.array(predictors)
model_linear = clf.fit(X, Y)
model_rbf_auto = clf_rbf_auto.fit(X,Y)
model_rbf_scale = clf_rbf_scale.fit(X,Y)

# Perform classification on samples in X.
y_prediction = clf.predict(X)

# Returns the mean accuracy on the given test data and labels.
linear_training_error = model_linear.score(X, Y)
print(linear_training_error)
#0.8516819571865444
rbf_score = model_rbf_auto.score(X,Y)
print(rbf_score)
#0.9938837920489296
rbf2_score = model_rbf_scale.score(X,Y)
print(rbf2_score)
#0.6605504587155964

# Weights assigned to the features
print('w = ',clf.coef_)

#Constant in the decision function
print('b = ',clf.intercept_)

####
print('Indices of support vectors = ', clf.support_)
print('Support vectors = ', clf.support_vectors_)
print('Number of support vectors for each class = ', clf.n_support_)
print('Coefficients of the support vector in the decision function = ', np.abs(clf.dual_coef_))

```

2. Using the k-nearest-neighbors classification function *kknn* contained in the R *kknn* package, suggest a good value of *k*, and show how well it classifies that data points in the full data set. Don't forget to scale the data (*scale=TRUE* in *kknn*).

Model1 has the best prediction accuracy with a value of 89.3%

```

> data = d0
> set.seed(9)
> rowidx = sample(1:nrow(data), round(.8*nrow(data)), replace=FALSE)
> X_train = data[rowidx,]; test = data[-rowidx,]
>
> library(kknn)
> model1 = kknn(class~, X_train, test, na.action=na.omit, k=11, distance=1, kernel='inv',scale=TRUE)
> model2 = kknn(class~, X_train, test, na.action=na.omit, k=11, distance=2, kernel='inv',scale=TRUE)

```

```
> model3 = kknnclass~, X_train, test, na.action=na.omit, k=11, distance=3, kernel='inv',scale=TRUE)
> model4 = kknnclass~, X_train, test, na.action=na.omit, k=11, distance=4, kernel='inv',scale=TRUE)
>
> model1accuracy = sum(model1$fitted.values == test[,11]) / length(test[,11])
> model2accuracy = sum(model2$fitted.values == test[,11]) / length(test[,11])
> model3accuracy = sum(model3$fitted.values == test[,11]) / length(test[,11])
> model4accuracy = sum(model4$fitted.values == test[,11]) / length(test[,11])
>
> model1accuracy
[1] 0.8931298
> model2accuracy
[1] 0.8854962
> model3accuracy
[1] 0.8778626
> model4accuracy
[1] 0.8778626
```